

01.03.2019

# Statistical Methods in AI (CSE/ECE 471)

Lecture-15: Bias-Variance, Model Selection, Regularization

Ravi Kiran

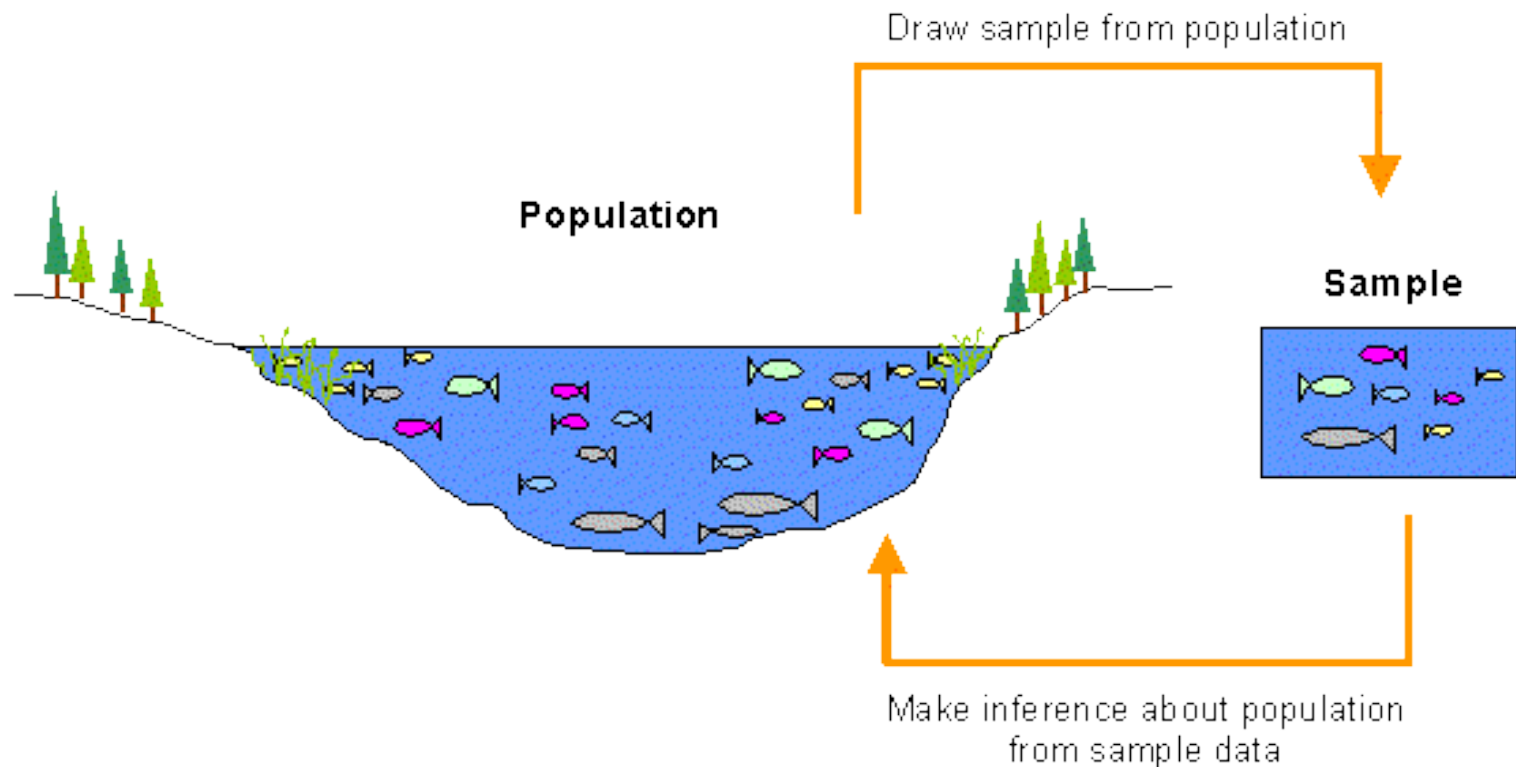
Center for Visual Information Technology (CVIT), IIIT Hyderabad



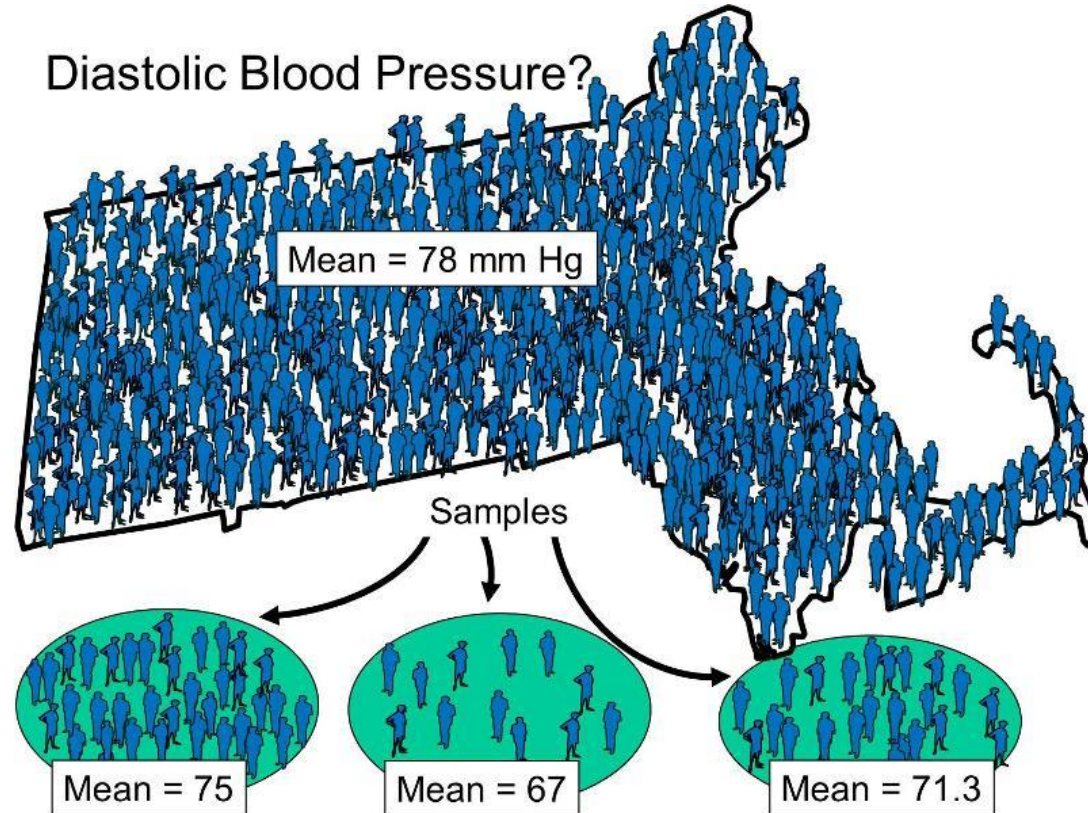
- ML workflow
  - Gather data
  - Select model (based on task)
  - Measure performance [using validation/test data]
  - Deploy model

- ML workflow
  - Gather data [“Data bias”]
  - Select model (based on task) [“Model bias”]
  - Measure performance [using validation/test data]
  - Deploy model

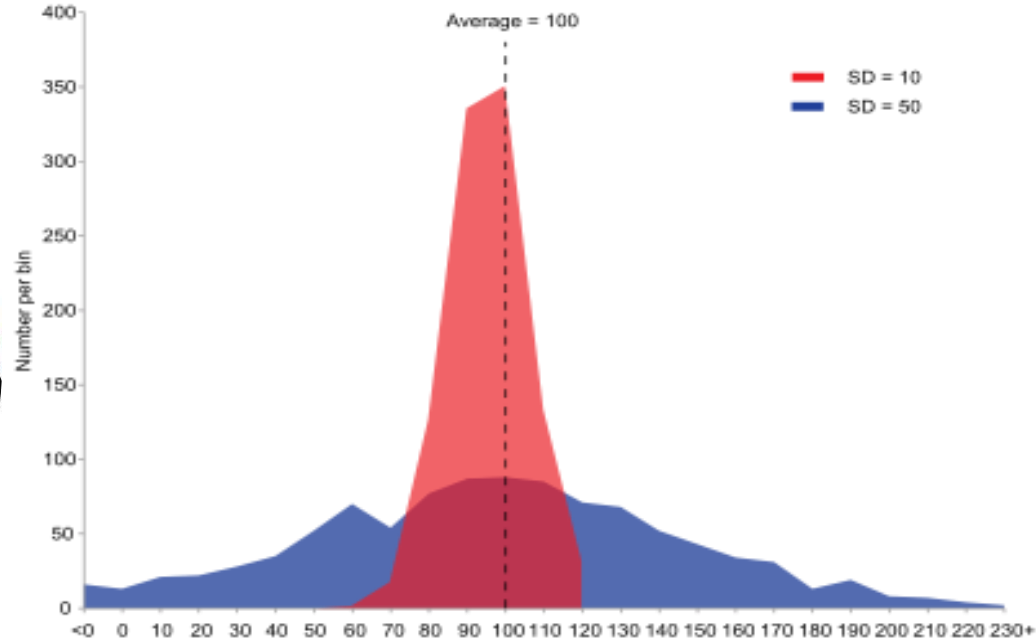
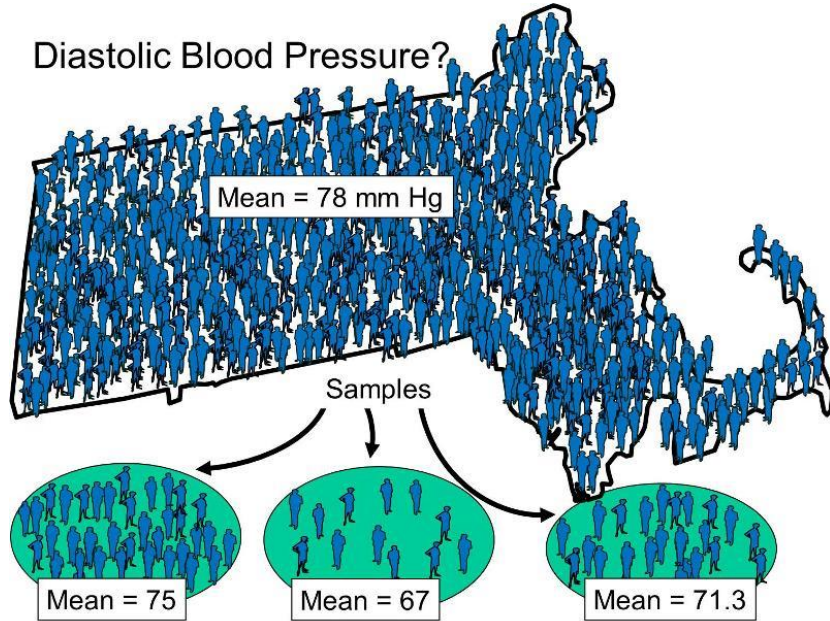
# Sample vs Population



# Statistics: Sample vs Population



# Statistics: Sample vs Population



- A good model should generalize well

# Data Collection Bias

- We place a huge amount of blind faith in the quality of datasets when training ML models.
- Quality
  - Coverage
    - Representative of Population
    - Label coverage
  - Noise
    - Input
    - Output [label noise]



# Data Representation Bias

- Input representation
- Output representation

# Model Bias

- Which 'type' of model ?
- Within the model 'type'
  - Model-specific choices

## Recall: Statistics 101

- Let  $X$  be a random variable with possible values  $x_i, i = 1 \dots n$  and with probability distribution  $P(X)$
- The *expected value* or *mean* of  $X$  is:

$$E[X] = \sum_{i=1}^n x_i P(x_i)$$

- If  $X$  is continuous, roughly speaking, the sum is replaced by an integral, and the distribution by a density function
- The *variance* of  $X$  is:

$$\begin{aligned} \text{Var}[X] &= E[(X - E(X))^2] \\ &= E[X^2] - (E[X])^2 \end{aligned}$$

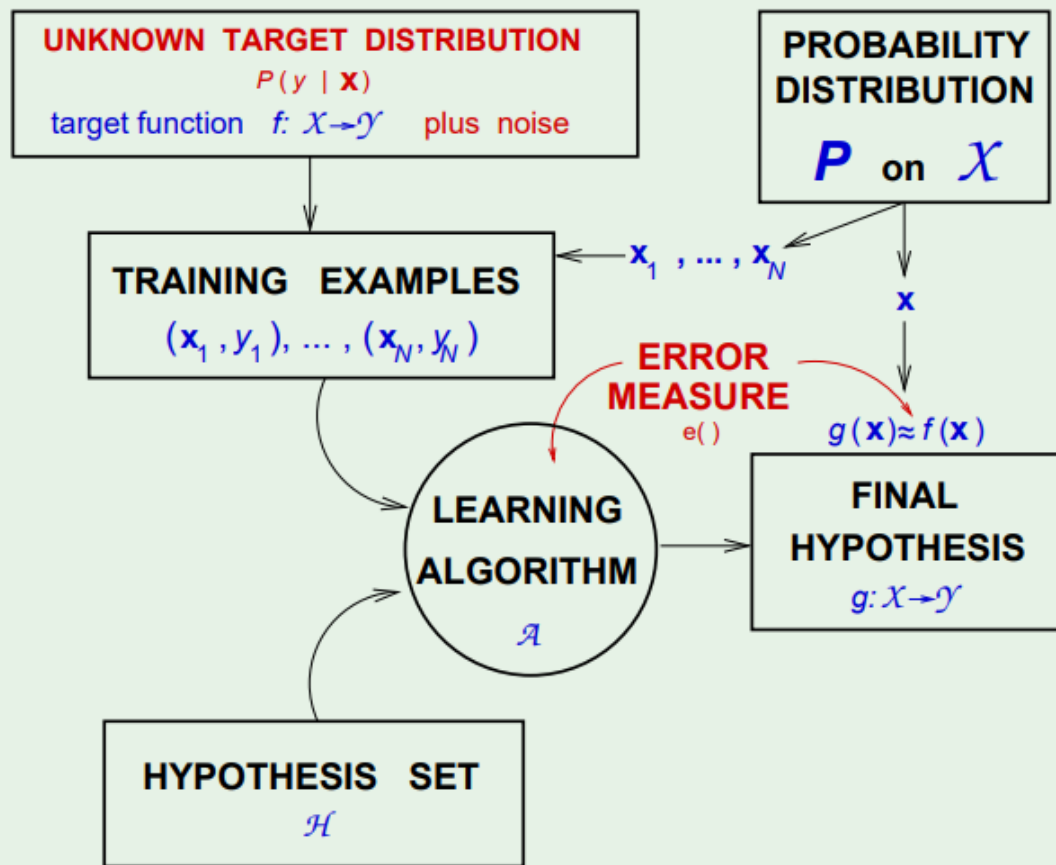
## The variance lemma

$$\begin{aligned} Var[X] &= E[(X - E[X])^2] \\ &= \sum_{i=1}^n (x_i - E[X])^2 P(x_i) \\ &= \sum_{i=1}^n (x_i^2 - 2x_i E[X] + (E[X])^2) P(x_i) \\ &= \sum_{i=1}^n x_i^2 P(x_i) - 2E[X] \sum_{i=1}^n x_i P(x_i) + (E[X])^2 \sum_{i=1}^n P(x_i) \\ &= E[X^2] - 2E[X]E[X] + (E[X])^2 \cdot 1 \\ &= E[X^2] - (E[X])^2 \end{aligned}$$

We will use the form:

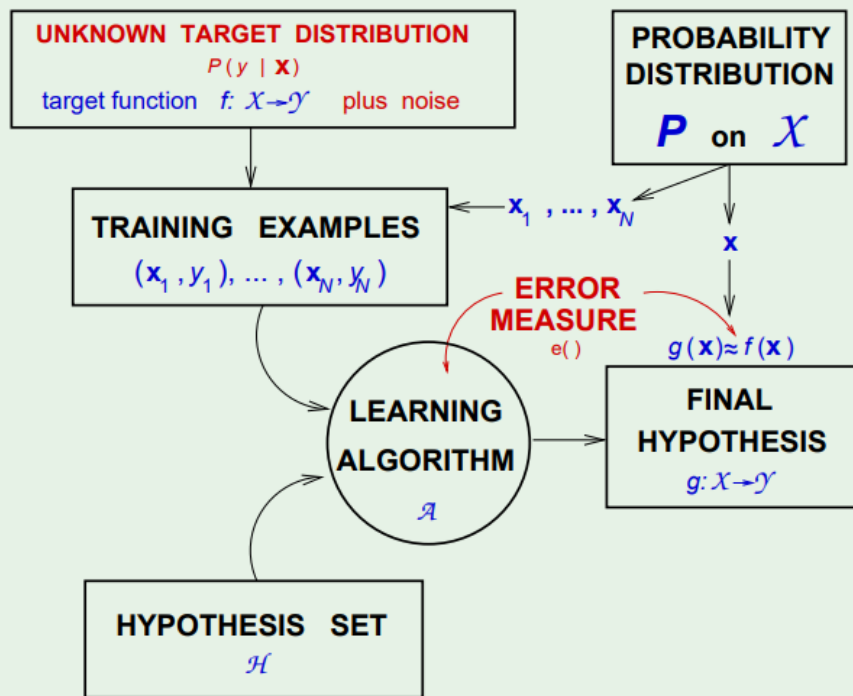
$$E[X^2] = (E[X])^2 + Var[X]$$

# The learning diagram - including noisy target



# Bias-Variance analysis

The learning diagram - including noisy target



1. How well  $\mathcal{H}$  can approximate  $f$

2. How well we can zoom in on a good  $h \in \mathcal{H}$

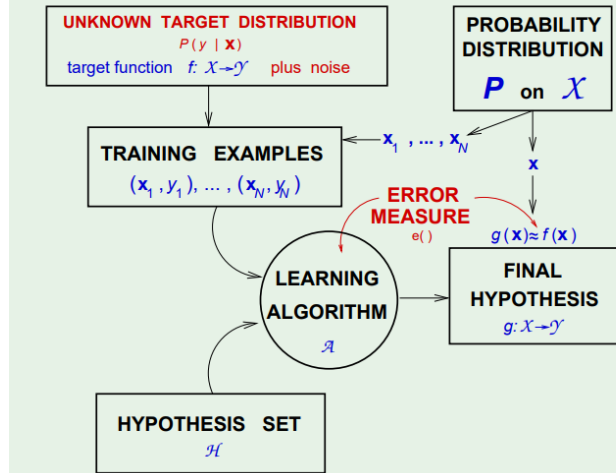
Applies to **real-valued targets** and uses **squared error**

$$E_{\text{out}}(g^{(\mathcal{D})}) = \left(g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x})\right)^2$$

$$E_{\text{out}}(g^{(\mathcal{D})}) = \mathbb{E}_{\mathbf{x}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right]$$

$$E_{\text{out}}(g^{(\mathcal{D})}) = \mathbb{E}_{\mathbf{x}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right]$$

The learning diagram - including noisy target





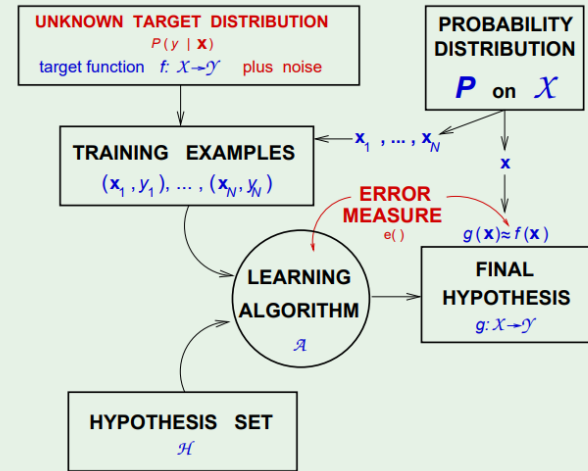
$$E_{\text{out}}(g^{(\mathcal{D})}) = \mathbb{E}_{\mathbf{x}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right]$$

$$E_{\text{out}}(g^{(\mathcal{D})}) = \mathbb{E}_{\mathbf{x}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right]$$

$$\mathbb{E}_{\mathcal{D}} [E_{\text{out}}(g^{(\mathcal{D})})] = \mathbb{E}_{\mathcal{D}} \left[ \mathbb{E}_{\mathbf{x}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right] \right]$$

$$= \mathbb{E}_{\mathbf{x}} \left[ \mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right] \right]$$

The learning diagram - including noisy target



# Bias-Variance analysis

$$\mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right]$$

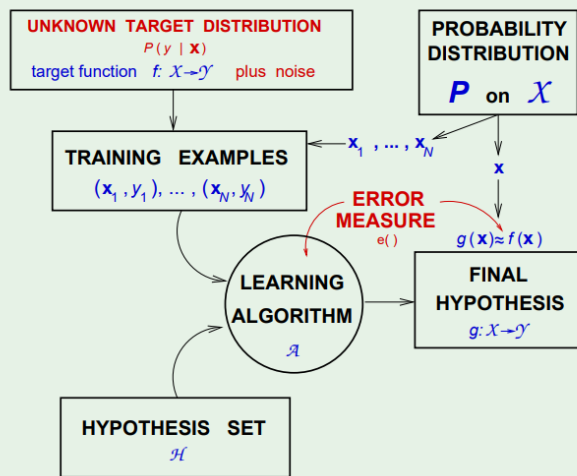
we define the 'average' hypothesis  $\bar{g}(\mathbf{x})$ :

$$\bar{g}(\mathbf{x}) = \mathbb{E}_{\mathcal{D}} \left[ g^{(\mathcal{D})}(\mathbf{x}) \right]$$

Imagine many data sets  $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_K$

$$\bar{g}(\mathbf{x}) \approx \frac{1}{K} \sum_{k=1}^K g^{(\mathcal{D}_k)}(\mathbf{x})$$

The learning diagram - including noisy target



# Bias-Variance analysis

$$\mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right] = \underbrace{\mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2 \right]}_{\text{var}(\mathbf{x})} + \underbrace{(\bar{g}(\mathbf{x}) - f(\mathbf{x}))^2}_{\text{bias}(\mathbf{x})}$$

$$\mathbb{E}_{\mathcal{D}} [E_{\text{out}}(g^{(\mathcal{D})})] = \mathbb{E}_{\mathbf{x}} \left[ \mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right] \right]$$

$$= \mathbb{E}_{\mathbf{x}} [\text{bias}(\mathbf{x}) + \text{var}(\mathbf{x})]$$

$$= \text{bias} + \text{var}$$

# Bias-Variance analysis

$$\mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right] = \underbrace{\mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2 \right]}_{\text{var}(\mathbf{x})} + \underbrace{(\bar{g}(\mathbf{x}) - f(\mathbf{x}))^2}_{\text{bias}(\mathbf{x})}$$

$$\mathbb{E}_{\mathcal{D}} [E_{\text{out}}(g^{(\mathcal{D})})] = \mathbb{E}_{\mathbf{x}} \left[ \mathbb{E}_{\mathcal{D}} \left[ (g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2 \right] \right]$$

$$= \mathbb{E}_{\mathbf{x}} [\text{bias}(\mathbf{x}) + \text{var}(\mathbf{x})]$$

$$= \text{bias} + \text{var}$$

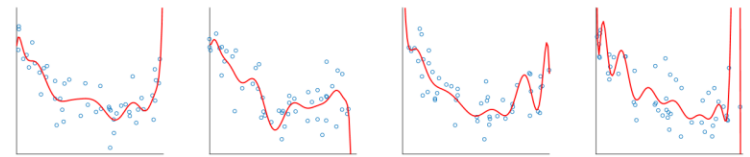
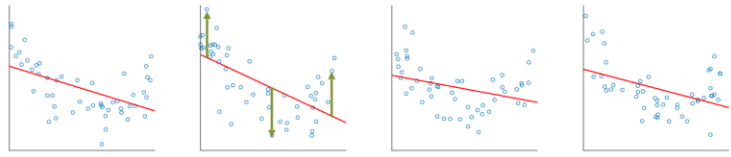
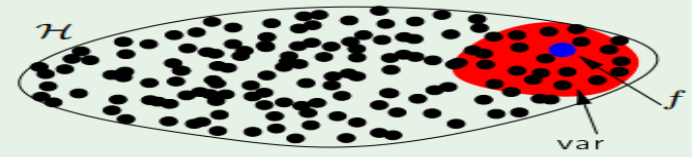
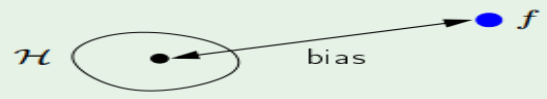
$$\mathbb{E}_{\mathcal{D}} \left[ \left( g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}) \right)^2 \right] = \underbrace{\mathbb{E}_{\mathcal{D}} \left[ \left( g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}) \right)^2 \right]}_{\text{var}(\mathbf{x})} + \underbrace{\left( \bar{g}(\mathbf{x}) - f(\mathbf{x}) \right)^2}_{\text{bias}(\mathbf{x})}$$

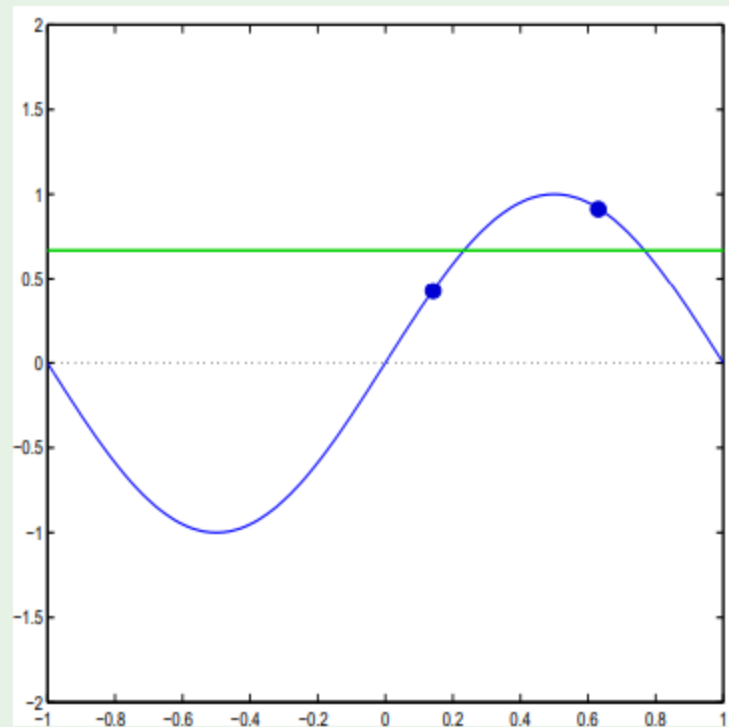
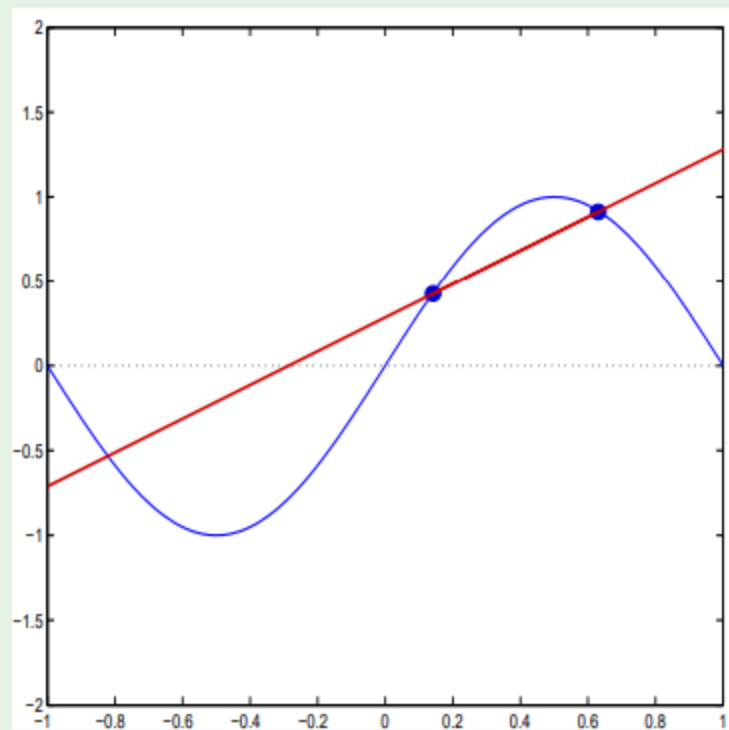
$$\mathbb{E}_{\mathbf{x}} [\text{bias}(\mathbf{x}) + \text{var}(\mathbf{x})]$$

The tradeoff

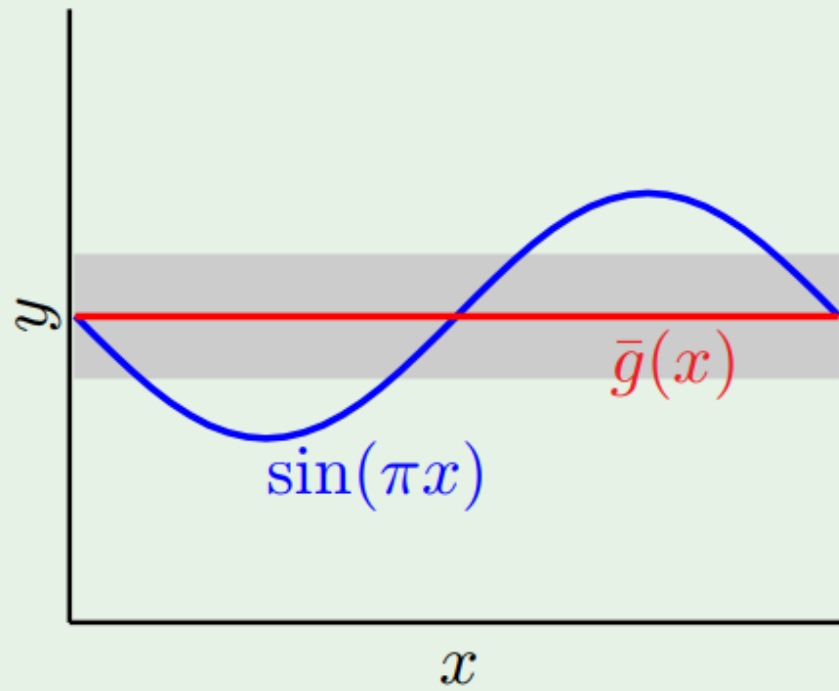
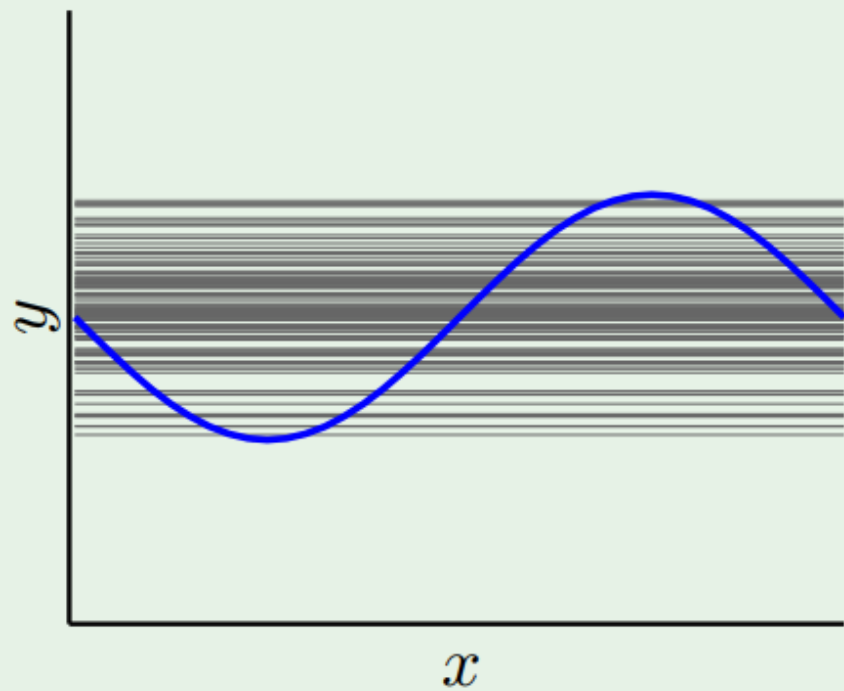
$$\text{bias} = \mathbb{E}_{\mathbf{x}} \left[ \left( \bar{g}(\mathbf{x}) - f(\mathbf{x}) \right)^2 \right]$$

$$\text{var} = \mathbb{E}_{\mathbf{x}} \left[ \mathbb{E}_{\mathcal{D}} \left[ \left( g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}) \right)^2 \right] \right]$$

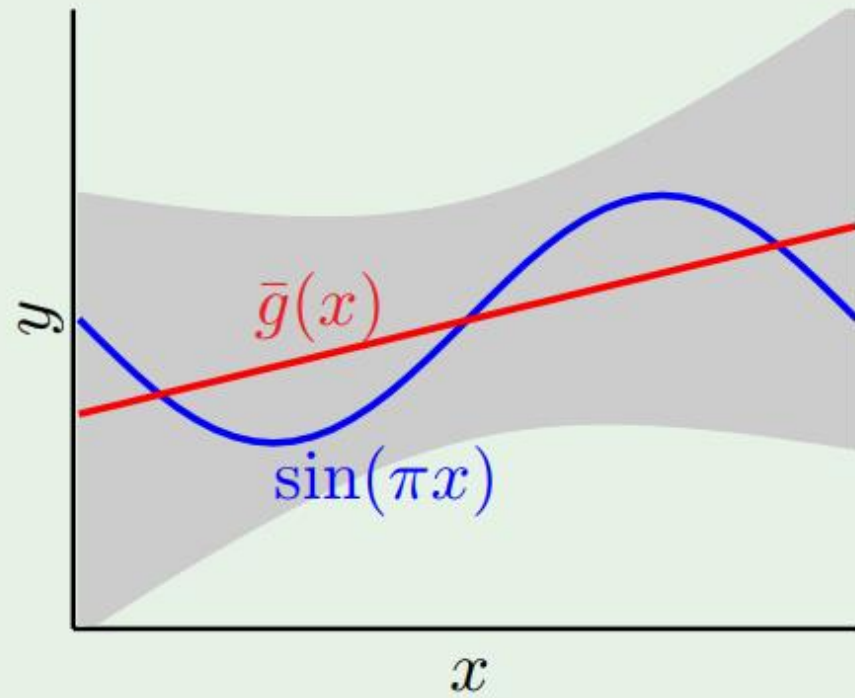
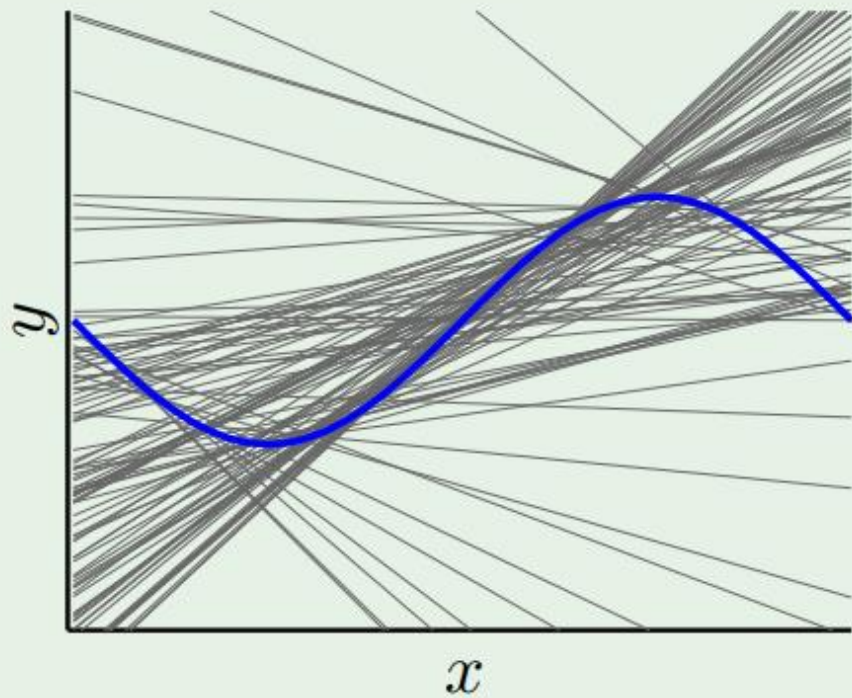


$\mathcal{H}_0$  $\mathcal{H}_1$ 

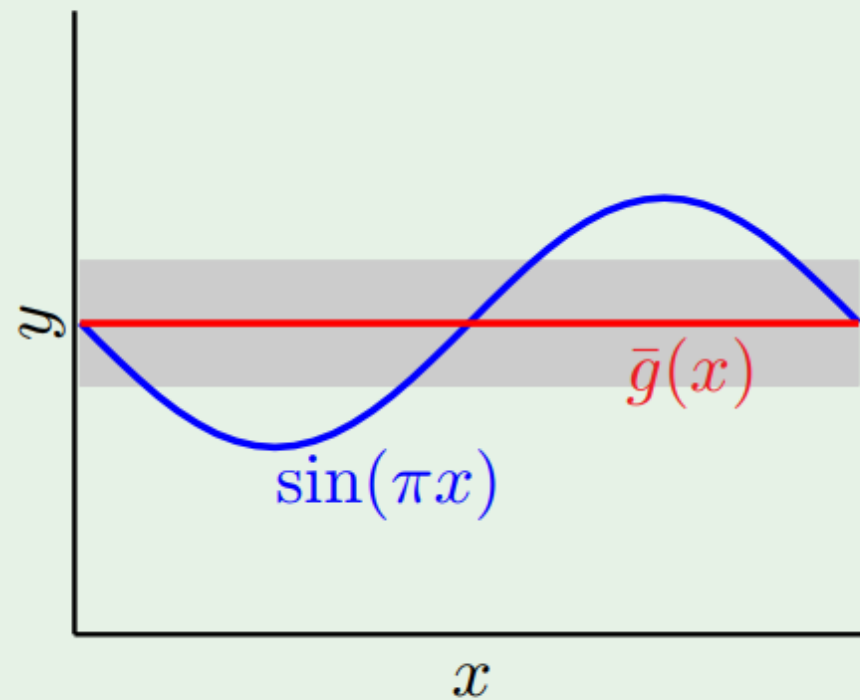
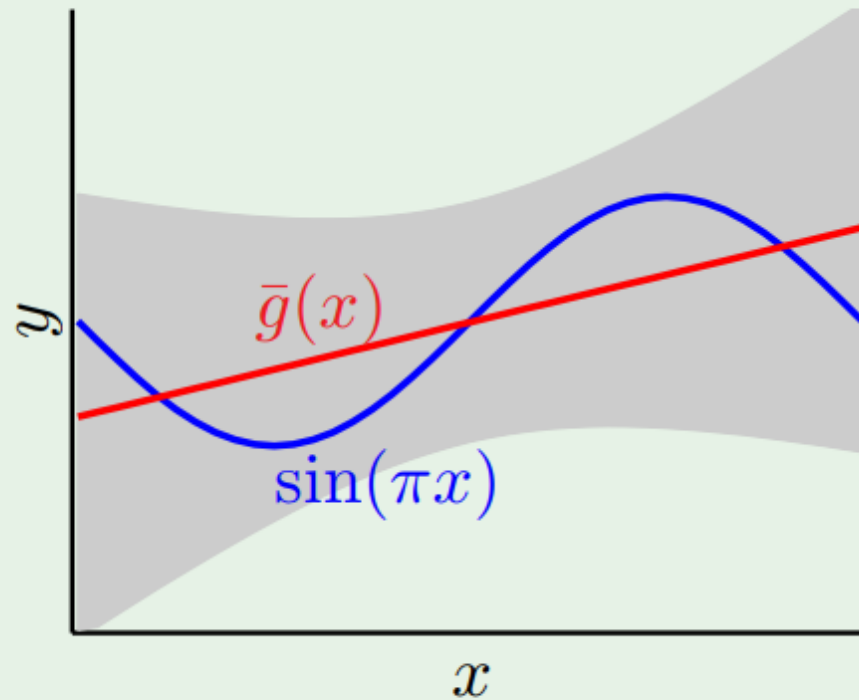
## Bias and variance - $\mathcal{H}_0$

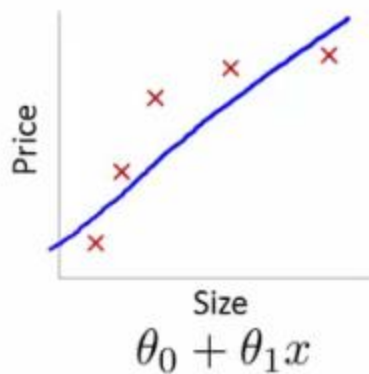


## Bias and variance - $\mathcal{H}_1$

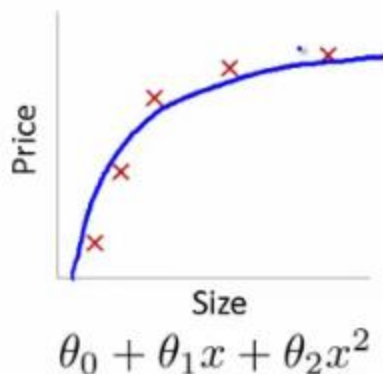




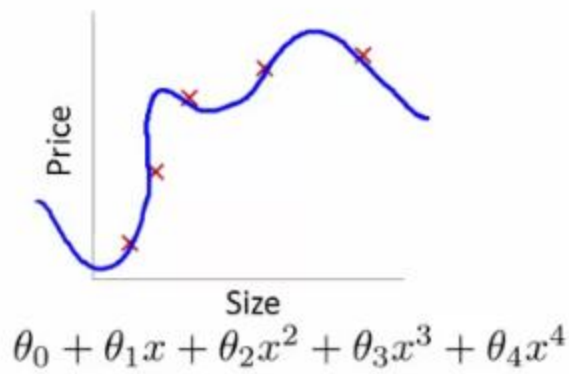
$\mathcal{H}_0$ bias = **0.50**var = **0.25** $\mathcal{H}_1$ bias = **0.21**var = **1.69**



High bias  
(underfit)



"Just right"



High variance  
(overfit)

- Bias-Variance decomposition provides insight into model complexity issue
- Limited practical value since it is based on ensembles of data sets
  - In practice there is only a single observed data set
  - If there are many training samples then combine them
    - which would reduce over-fitting for a given model complexity

## Regularization

- Remember the intuition: complicated hypotheses lead to overfitting
- Idea: change the error function to *penalize hypothesis complexity*:

$$J(\mathbf{w}) = J_D(\mathbf{w}) + \lambda J_{pen}(\mathbf{w})$$

This is called *regularization* in machine learning and *shrinkage* in statistics

- $\lambda$  is called *regularization coefficient* and controls how much we value fitting the data well, vs. a simple hypothesis

## Regularization for linear models

- A squared penalty on the weights would make the math work nicely in our case:

$$\frac{1}{2}(\Phi \mathbf{w} - \mathbf{y})^T(\Phi \mathbf{w} - \mathbf{y}) + \frac{\lambda}{2}\mathbf{w}^T\mathbf{w}$$

- This is also known as  *$L_2$  regularization*, or *weight decay* in neural networks
- By re-grouping terms, we get:

$$J_D(\mathbf{w}) = \frac{1}{2}(\mathbf{w}^T(\Phi^T\Phi + \lambda\mathbf{I})\mathbf{w} - \mathbf{w}^T\Phi^T\mathbf{y} - \mathbf{y}^T\Phi\mathbf{w} + \mathbf{y}^T\mathbf{y})$$

- Optimal solution (obtained by solving  $\nabla_{\mathbf{w}}J_D(\mathbf{w}) = 0$ )

$$\mathbf{w} = (\Phi^T\Phi + \lambda\mathbf{I})^{-1}\Phi^T\mathbf{y}$$

## What $L_2$ regularization does

$$\arg \min_{\mathbf{w}} \frac{1}{2}(\Phi \mathbf{w} - \mathbf{y})^T(\Phi \mathbf{w} - \mathbf{y}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} = (\Phi^T \Phi + \lambda I)^{-1} \Phi^T \mathbf{y}$$

- If  $\lambda = 0$ , the solution is the same as in regular least-squares linear regression
- If  $\lambda \rightarrow \infty$ , the solution  $\mathbf{w} \rightarrow 0$
- Positive  $\lambda$  will cause the magnitude of the weights to be smaller than in the usual linear solution
- This is also called *ridge regression*, and it is a special case of Tikhonov regularization
- A different view of regularization: we want to optimize the error while keeping the  $L_2$  norm of the weights,  $\mathbf{w}^T \mathbf{w}$ , bounded.

## $L_2$ Regularization for linear models revisited

- Optimization problem: minimize error while keeping norm of the weights bounded

$$\begin{aligned} \min_{\mathbf{w}} J_D(\mathbf{w}) &= \min_{\mathbf{w}} (\Phi \mathbf{w} - \mathbf{y})^T (\Phi \mathbf{w} - \mathbf{y}) \\ \text{such that } \mathbf{w}^T \mathbf{w} &\leq \eta \end{aligned}$$

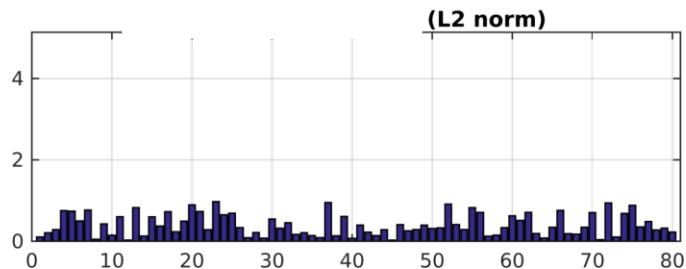
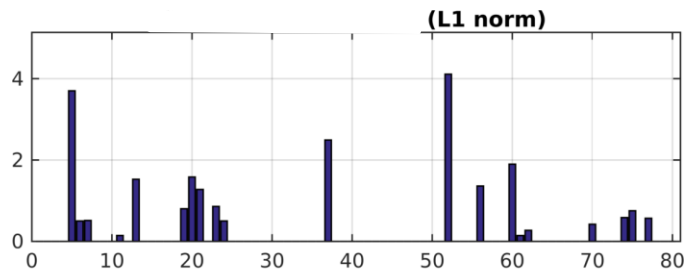
- The Lagrangian is:

$$L(\mathbf{w}, \lambda) = J_D(\mathbf{w}) - \lambda(\eta - \mathbf{w}^T \mathbf{w}) = (\Phi \mathbf{w} - \mathbf{y})^T (\Phi \mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w} - \lambda \eta$$

- For a fixed  $\lambda$ , and  $\eta = \lambda^{-1}$ , the best  $\mathbf{w}$  is the same as obtained by weight decay

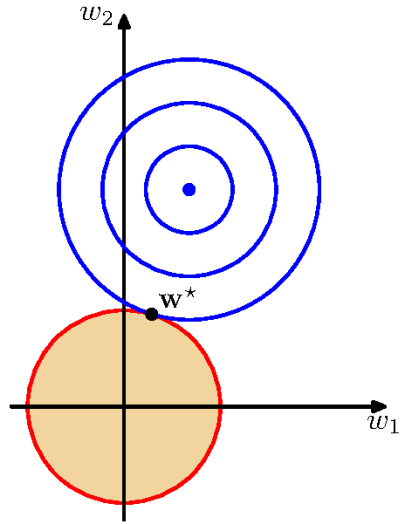
## Pros and cons of $L_2$ regularization

- If  $\lambda$  is at a “good” value, regularization helps to avoid overfitting
- Choosing  $\lambda$  may be hard: cross-validation is often used
- If there are irrelevant features in the input (i.e. features that do not affect the output),  $L_2$  will give them small, but non-zero weights.
- Ideally, irrelevant input should have weights exactly equal to 0.





## Visualizing regularization (2 parameters)



$$\mathbf{w}^* = (\Phi^T \Phi + \lambda I)^{-1} \Phi \mathbf{y}$$

## $L_1$ Regularization for linear models

- Instead of requiring the  $L_2$  norm of the weight vector to be bounded, make the requirement on the  $L_1$  norm:

$$\min_{\mathbf{w}} J_D(\mathbf{w}) = \min_{\mathbf{w}} (\Phi \mathbf{w} - \mathbf{y})^T (\Phi \mathbf{w} - \mathbf{y})$$

such that  $\sum_{i=1}^n |w_i| \leq \eta$

- This yields an algorithm called Lasso (Tibshirani, 1996)

## Solving $L_1$ regularization

- The optimization problem is a quadratic program
- There is one constraint for each possible sign of the weights ( $2^n$  constraints for  $n$  weights)
- For example, with two weights:

$$\min_{w_1, w_2} \sum_{j=1}^m (y_j - w_1 x_1 - w_2 x_2)^2$$

$$\text{such that } w_1 + w_2 \leq \eta$$

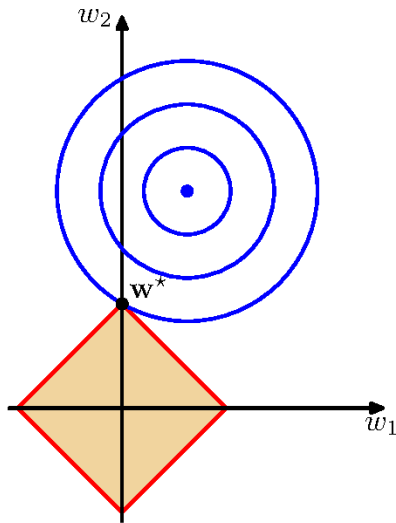
$$w_1 - w_2 \leq \eta$$

$$-w_1 + w_2 \leq \eta$$

$$-w_1 - w_2 \leq \eta$$

- Solving this program directly can be done for problems with a small number of inputs

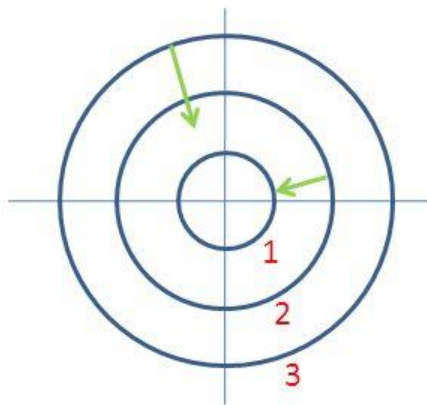
## Visualizing $L_1$ regularization



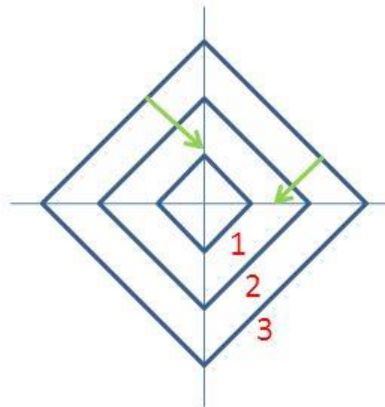
- If  $\lambda$  is big enough, the circle is very likely to intersect the diamond at one of the corners
- This makes  $L_1$  regularization much more likely to make some weights *exactly* 0

# A nice property of $L_1$

- Gradients of  $L_2$ - and  $L_1$ -norm



Negative gradient of  $L_2$ -norm  
always points directly toward 0

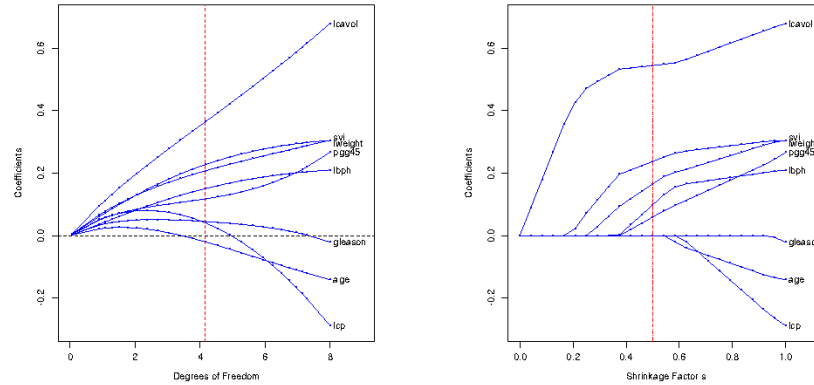


"Negative gradient" of  $L_1$ -norm  
(direction of steepest descent)  
points toward coordinate axes

## Pros and cons of $L_1$ regularization

- If there are irrelevant input features, Lasso is likely to make their weights 0, while  $L_2$  is likely to just make all weights small
- Lasso is biased towards providing *sparse solutions* in general
- Lasso optimization is computationally more expensive than  $L_2$
- More efficient solution methods have to be used for large numbers of inputs (e.g. least-angle regression, 2003).
- $L_1$  methods of various types are very popular

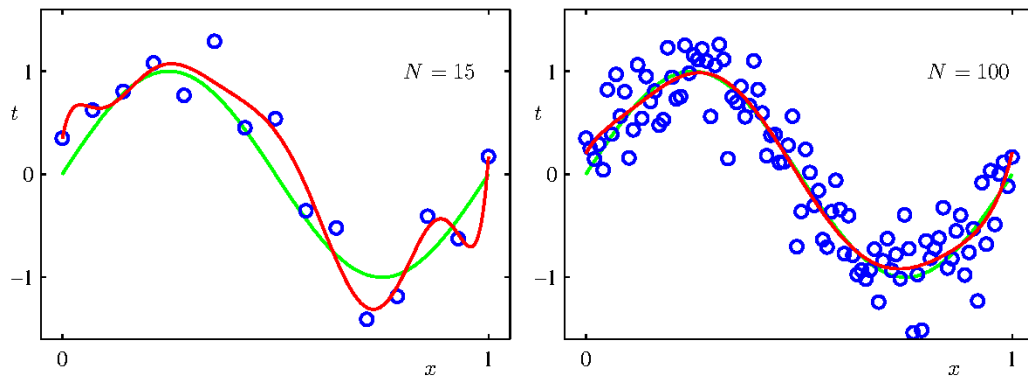
## Example of L1 vs L2 effect



- Note the sparsity in the coefficients induced by  $L_1$
- Lasso is an efficient way of performing the  $L_1$  optimization

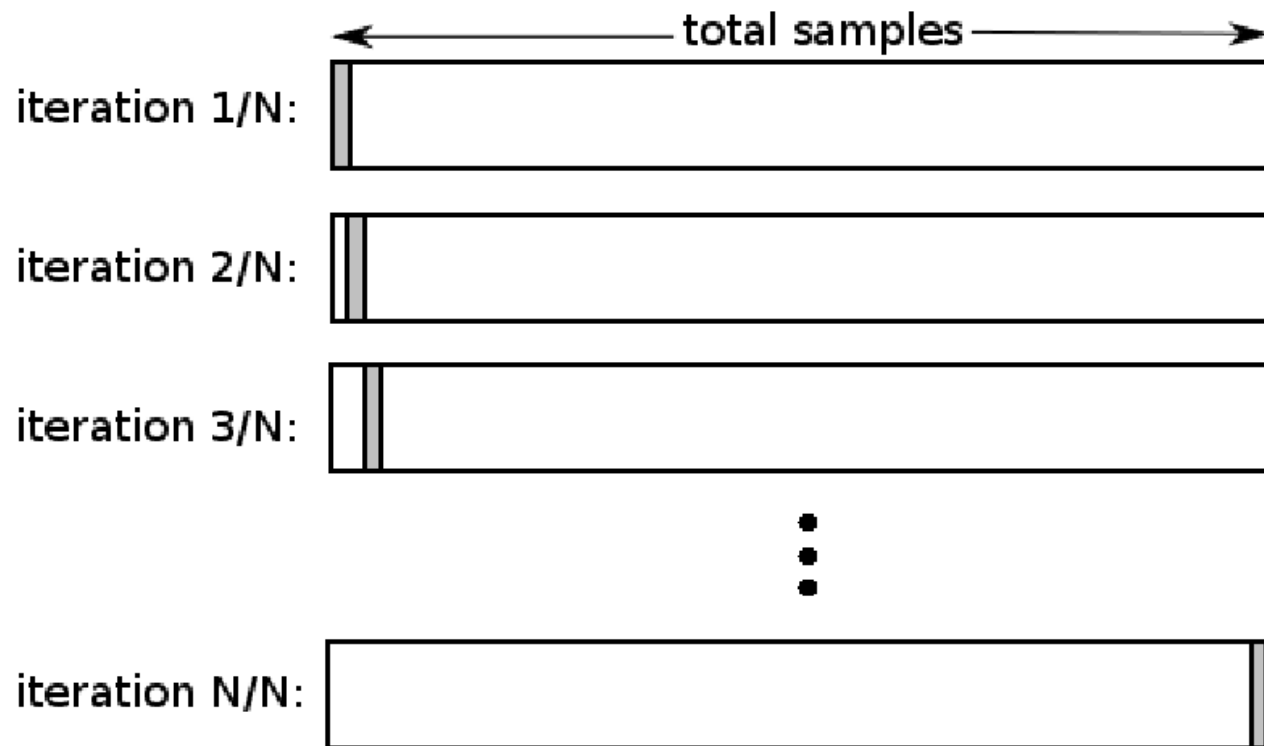
## More on overfitting

- Overfitting depends on the amount of data, relative to the complexity of the hypothesis
- With more data, we can explore more complex hypotheses spaces, and still find a good solution



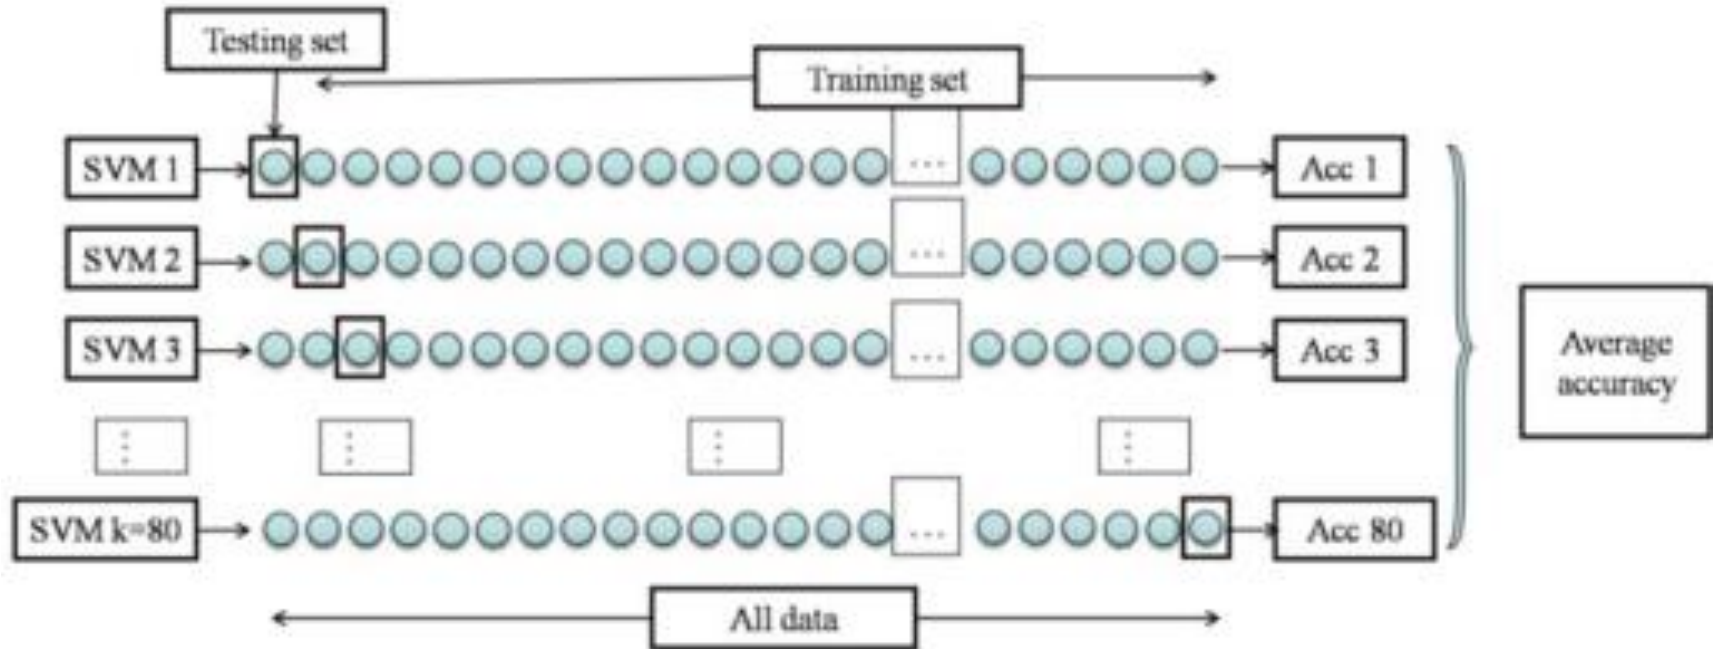


## Leave One Out Cross Validation



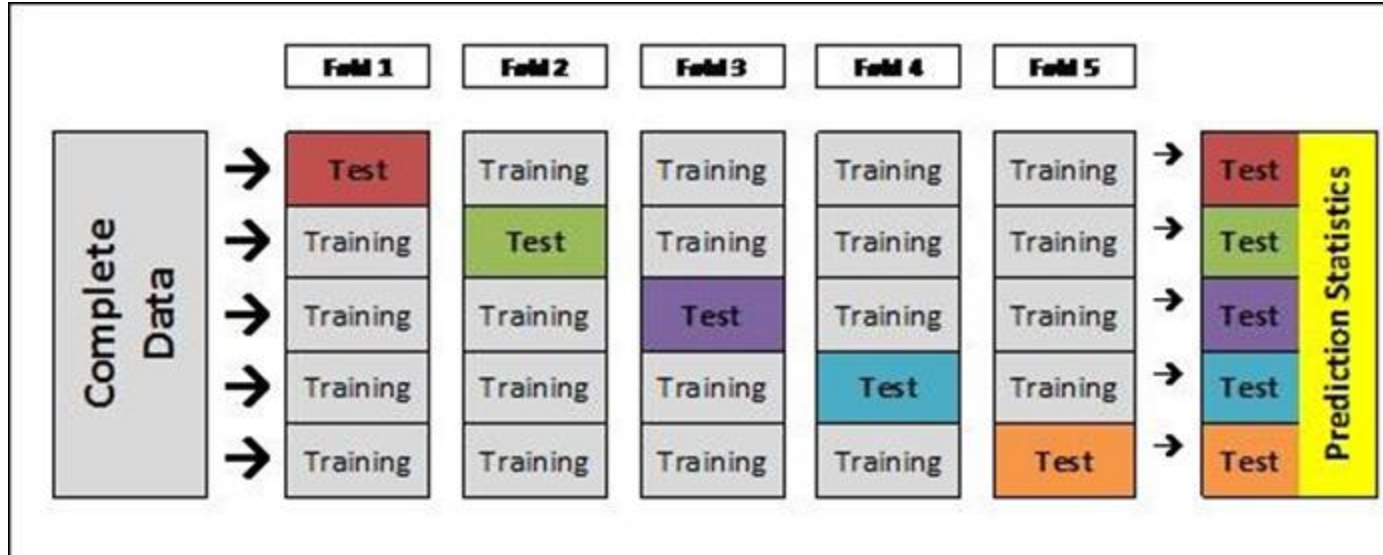
## Leave One Out Cross Validation

**High variance**



## K-fold Cross Validation

Fold → bias



*Assess model's bias-variance characteristics using k-fold Cross Validation?*

- With  $k$ -fold CV , you get  $k$  different estimates of your model's error-  $e_1, e_2, e_3, \dots, e_k$
- $\text{Mean}(\text{errors}) \sim 0 \rightarrow \text{low bias}$
- $\text{Variance}(\text{errors}) \sim 0 \rightarrow \text{low variance}$