

22.03.2019

Statistical Methods in AI (CSE/ECE 471)

Lecture-19: ML for Sequential Data - Hidden Markov Models

~

Ravi Kiran

Center for Visual Information Technology (CVIT), IIIT Hyderabad



Announcements


- For final exam
 - Any answer written with a pencil will AUTOMATICALLY get 0 marks !

Analysis of Sequential Data

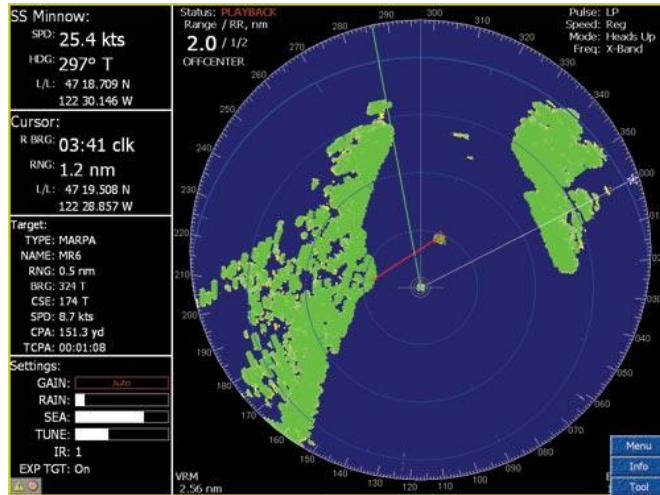
- Sequential structure arises in a huge range of applications
 - Repeated measurements of a temporal process
 - Online decision making & control
 - Text, biological sequences etc.

Speech Recognition





- Given an audio waveform, robustly extract & recognize any spoken words
- Statistical models can be used to
 - Provide greater robustness to noise
 - Adapt to accent of different speakers
 - Learn from training



Target Tracking



*Radar-based tracking
of multiple targets*




*Visual tracking of
articulated objects*
(L. Sigal et. al., 2006)

- Estimate motion of targets in 3D world from indirect, potentially noisy measurements


Robot Navigation: *SLAM*

Simultaneous Localization and Mapping




*CAD
Map*

(S. Thrun,
San Jose Tech Museum)



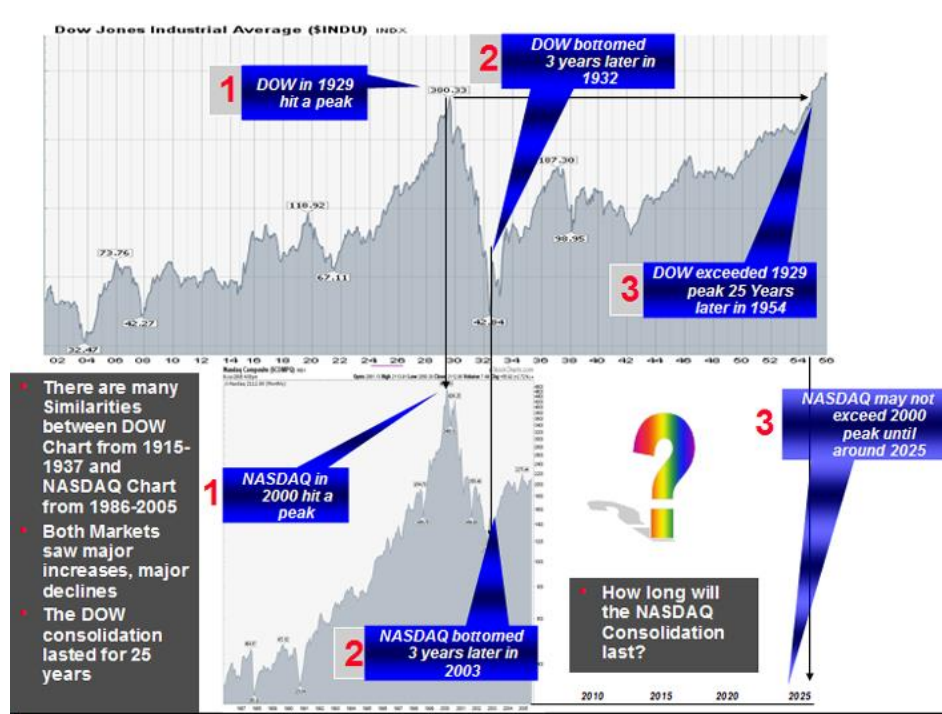
*Estimated
Map*

*Landmark
SLAM
(E. Nebot,
Victoria Park)*



- As robot moves, estimate its pose & world geometry

Financial Forecasting



<http://www.steadfastinvestor.com/>

- Predict future market behavior from historical data, news reports, expert opinions, ...

i.i.d to sequential data


- ❑ So far we assumed independent, identically distributed data

$$\{X_i\}_{i=1}^n \stackrel{iid}{\sim} p(\mathbf{X})$$

- ❑ Sequential (non i.i.d.) data

- Time-series data

- E.g. Speech



- Characters in a sentence



- Base pairs along a DNA strand



Sequential Processes

- Consider a system which can occupy one of N discrete *states* or *categories*

$$x_t \in \{1, 2, \dots, N\} \rightarrow \text{state at time } t$$

- We are interested in *stochastic* systems, in which state evolution is random
- Any *joint* distribution can be factored into a series of *conditional* distributions:

$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t \mid x_0, \dots, x_{t-1})$$

$$\begin{aligned} p(A, B) &= p(A) p(B|A) \\ p(A, B, C) &= p(A) p(B|A) p(C|A, B) \end{aligned}$$

Markov Processes

- For a *Markov* process, the next state depends only on the current state:

$$p(x_{t+1} \mid x_0, \dots, x_t) = p(x_{t+1} \mid x_t)$$

- This property in turn implies that

$$\begin{aligned} & p(x_0, \dots, x_{t-1}, x_{t+1}, \dots, x_T \mid x_t) \\ &= p(x_0, \dots, x_{t-1} \mid x_t) p(x_{t+1}, \dots, x_T \mid x_t) \end{aligned}$$

*“Conditioned on the present,
the past & future are independent”*

State Transition Matrices

x_1 x_2 x_3 $x_1 = a_3$ $x_2 = a_5$ $x_3 = a_2 \dots$

N • A *stationary* Markov chain with N states is described by an $N \times N$ *transition matrix*:

$x_1 = a_1$
 \vdots
 a_N
 a_3 $a_5 \dots$

$$Q = \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix}$$

$$q_{ij} \triangleq p(x_{t+1} = i \mid x_t = j)$$


- Constraints on valid transition matrices:

$$q_{ij} \geq 0 \quad \sum_{i=1}^N q_{ij} = 1 \quad \text{for all } j$$

State Transition Diagrams

$$q_{ij} \triangleq p(x_{t+1} = i \mid x_t = j)$$

$$Q = \begin{bmatrix} 0.5 & 0.1 & 0.0 \\ 0.3 & 0.0 & 0.4 \\ 0.2 & 0.9 & 0.6 \end{bmatrix}$$



- Think of a particle randomly following an arrow at each discrete time step
- Most useful when N small, and Q *sparse*

Markov Models

□ Markov Assumption

parameters in
stationary model
K-ary variables

$$1^{\text{st}} \text{ order} \quad p(\mathbf{X}) = \prod_{i=1}^n p(X_i | X_{i-1}) \quad O(K^2)$$

$$m^{\text{th}} \text{ order} \quad p(\mathbf{X}) = \prod_{i=1}^n p(X_i | X_{i-1}, \dots, X_{i-m}) \quad O(K^{m+1})$$

$$n-1^{\text{th}} \text{ order} \quad p(\mathbf{X}) = \prod_{i=1}^n p(X_i | X_{i-1}, \dots, X_1) \quad O(K^n)$$

≡ no assumptions – complete (but directed) graph


Homogeneous/stationary Markov model (probabilities don't depend on n)

Hidden Markov Models

- Few realistic time series directly satisfy the assumptions of Markov processes:


*“Conditioned on the present,
the past & future are independent”*

- Motivates *hidden Markov models (HMM)*:



$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t | x_{t-1}) p(y_t | x_t)$$


Hidden states



- Given x_t , earlier observations provide no *additional information* about the future:


$$p(y_t, y_{t+1}, \dots \mid x_t, y_{t-1}, y_{t-2}, \dots) = p(y_t, y_{t+1}, \dots \mid x_t)$$

Where do states come from?



- Analysis of a *physical phenomenon*:
 - Dynamical models of an aircraft or robot
 - Geophysical models of climate evolution
- Discovered from *training data*:
 - Recorded examples of spoken English
 - Historic behavior of stock prices


Discrete State HMMs



- Associate each of the N hidden states with a different observation distribution:

$$p(y_t | x_t = 1) \quad p(y_t | x_t = 2) \quad \dots$$

- Observation densities are typically chosen to encode domain knowledge



Discrete HMMs: Observations

Discrete Observations



$$y_t \in \{1, 2, \dots, M\}$$

$$p(y_t \mid x_t = 1) = \begin{bmatrix} 0.3 \\ 0.1 \\ 0.5 \\ 0.1 \end{bmatrix}$$


$$p(y_t \mid x_t = 2) = \begin{bmatrix} 0.2 \\ 0.2 \\ 0.1 \\ 0.5 \end{bmatrix}$$

Continuous Observations

$$y_t \in \mathbb{R}^k$$




Specifying an HMM



- Observation model: $P(y_i|x_i)$
- Transition model: $P(x_i|x_{i-1})$
- Initial state distribution: $P(x_0)$

Hidden Markov Models



$$p(S_1, \dots, S_T, O_1, \dots, O_T) = \prod_{t=1}^T p(O_t | S_t) \prod_{t=1}^T p(S_t | S_{t-1})$$

1st order Markov assumption on hidden states $\{S_t\} \ t = 1, \dots, T$
(can be extended to higher order).

Note: O_t depends on all previous observations $\{O_{t-1}, \dots, O_1\}$

$$\begin{aligned} p(A, B) \\ p(A|B) &= \frac{p(A, B)}{p(B)} \\ &= \frac{p(A, B)}{\sum_A p(A, B)} \end{aligned}$$

Hidden Markov Models

- Parameters – stationary/homogeneous markov model (independent of time t)

Initial probabilities


$$p(S_1 = i) = \pi_i$$

Transition probabilities

$$p(S_t = j | S_{t-1} = i) = p_{ij}$$

Emission probabilities

$$p(O_t = y | S_t = i) = q_i^y$$



$$p(\{S_t\}_{t=1}^T, \{O_t\}_{t=1}^T) =$$

$$p(S_1) \prod_{t=2}^T p(S_t | S_{t-1}) \prod_{t=1}^T p(O_t | S_t)$$

HMM Example

- The Dishonest Casino

A casino has two dices:

Fair dice

$$P(1) = P(2) = P(3) = P(5) = P(6) = 1/6$$

Loaded dice

$$P(1) = P(2) = P(3) = P(5) = 1/10$$

$$P(6) = \frac{1}{2}$$

Casino player switches back-&-forth between fair and loaded die with 5% probability



HMM Problems

GIVEN: A sequence of rolls by the casino player


1245526462146146136136661664661636616366163616515615115146123562344

QUESTION

- How likely is this sequence, given our model of how the casino works?
 - This is the **EVALUATION** problem in HMMs
- What portion of the sequence was generated with the fair die, and what portion with the loaded die?
 - This is the **DECODING** question in HMMs
- How “loaded” is the loaded die? How “fair” is the fair die? How often does the casino player change from fair to loaded, and back?
 - This is the **LEARNING** question in HMMs

HMM Example

- Observed sequence: $\{O_t\}_{t=1}^T$




- Hidden sequence $\{S_t\}_{t=1}^T$ (or segmentation):



State Space Representation

- Switch between **F** and **L** with 5% probability



- HMM Parameters**

Initial probs

$$P(S_1 = \text{L}) = 0.5 = P(S_1 = \text{F})$$

Transition probs

$$P(S_t = \text{L/F} | S_{t-1} = \text{L/F}) = 0.95$$

$$P(S_t = \text{F/L} | S_{t-1} = \text{L/F}) = 0.05$$

Emission probabilities

$$P(O_t = y | S_t = \text{F}) = 1/6 \quad y = 1, 2, 3, 4, 5, 6$$

$$P(O_t = y | S_t = \text{L}) = 1/10 \quad y = 1, 2, 3, 4, 5$$
$$= 1/2 \quad y = 6$$

Three main problems in HMMs

- 
- **Evaluation** – Given HMM parameters & observation seqn $\{O_t\}_{t=1}^T$

find $p(\{O_t\}_{t=1}^T | \theta)$ prob of observed sequence

- How likely is this sequence, given our model of how the casino works?

- **Decoding** – Given HMM parameters & observation seqn $\{O_t\}_{t=1}^T$

find $\arg \max_{s_1, \dots, s_T} p(\{S_t\}_{t=1}^T | \{O_t\}_{t=1}^T, \theta)$ most probable

sequence of hidden states

- What portion of the sequence was generated with the fair die, and what portion with the loaded die?

- **Learning** – Given HMM with unknown parameters and $\{O_t\}_{t=1}^T$ observation sequence

find $\arg \max_{\theta} p(\{O_t\}_{t=1}^T | \theta)$ parameters that maximize

likelihood of observed data

- How “loaded” is the loaded die? How “fair” is the fair die? How often does the casino player change from fair to loaded, and back?

HMM Algorithms

- **Evaluation** – What is the probability of the observed sequence? **Forward Algorithm**
- **Decoding** – What is the probability that the third roll was loaded given the observed sequence? **Forward-Backward Algorithm**
 - What is the most likely die sequence given the observed sequence? **Viterbi Algorithm**
- **Learning** – Under what parameterization is the observed sequence most probable? **Baum-Welch Algorithm (EM)**


Evaluation Problem

- Given HMM parameters $p(S_1), p(S_t|S_{t-1}), p(O_t|S_t)$ & observation sequence $\{O_t\}_{t=1}^T$

find probability of observed sequence

$$\underline{p(\{O_t\}_{t=1}^T)} = \sum_{S_1, \dots, S_T} \underline{p(\{O_t\}_{t=1}^T, \{S_t\}_{t=1}^T)}$$

$$= \sum_{S_1, \dots, S_T} p(S_1) \prod_{t=2}^T p(S_t|S_{t-1}) \prod_{t=1}^T p(O_t|S_t)$$



$$\sum_{S_1} \sum_{S_2} \dots \sum_{S_T}$$

requires summing over all possible hidden state values at all times – K^T exponential # terms!

Instead: $p(\{O_t\}_{t=1}^T) = \sum_k \underbrace{p(\{O_t\}_{t=1}^T, S_T = k)}$

α_T^k Compute recursively


Forward Probability

$$p(\{O_t\}_{t=1}^T) = \sum_k p(\{O_t\}_{t=1}^T, S_T = k) = \sum_k \alpha_T^k$$

Compute forward probability α_t^k recursively over t

$$\alpha_t^k := p(O_1, \dots, O_t, S_t = k)$$

Handwritten notes: $t \leq T$ (underlined), $S_{t-1} = i$ (above the ellipsis)




Introduce S_{t-1}

Chain rule

Markov assumption

$$= p(O_t | S_t = k) \sum_{i=1}^K \alpha_{t-1}^i p(S_t = k | S_{t-1} = i)$$

Handwritten notes: K (above the sum), α_{t-1}^i (circled), $S_{t-1} = i$ (underlined)



Forward Algorithm

Can compute α_t^k for all k, t using dynamic programming:

- Initialize: $\alpha_1^k = p(O_1 | S_1 = k) p(S_1 = k)$ for all k

- Iterate: for $t = 2, \dots, T$

$$\alpha_t^k = p(O_t | S_t = k) \sum_i \alpha_{t-1}^i p(S_t = k | S_{t-1} = i) \quad \text{for all } k$$

- Termination: $p(\{O_t\}_{t=1}^T) = \sum_k \alpha_T^k$